

1.) Import the Credit Card Fraud Data From CCLE

```
import pandas as pd
from google.colab import drive
import matplotlib.pyplot as plt
import numpy as np

drive.mount('/content/gdrive/', force_remount = True)

Mounted at /content/gdrive/

df = pd.read_csv("/content/gdrive/MyDrive/Econ 441B/fraudTest.csv")
```

df.head()



|   | Unnamed: 0 | trans_date_trans_time | cc_num           | merchant                             | category       | amt   | first  | last     | gender | street                      | ... | lat     | long      | city_pop | job                    | dob        |              |
|---|------------|-----------------------|------------------|--------------------------------------|----------------|-------|--------|----------|--------|-----------------------------|-----|---------|-----------|----------|------------------------|------------|--------------|
| 0 | 0          | 2020-06-21 12:14:25   | 2291163933867244 | fraud_Kirlin and Sons                | personal_care  | 2.86  | Jeff   | Elliott  | M      | 351 Darlene Green           | ... | 33.9659 | -80.9355  | 333497   | Mechanical engineer    | 1968-03-19 | 2da90c7d74bc |
| 1 | 1          | 2020-06-21 12:14:33   | 3573030041201292 | fraud_Sporer-Keebler                 | personal_care  | 29.84 | Joanne | Williams | F      | 3638 Marsh Union            | ... | 40.3207 | -110.4360 | 302      | Sales professional, IT | 1990-01-17 | 324cc2044076 |
| 2 | 2          | 2020-06-21 12:14:53   | 3598215285024754 | fraud_Swaniawski, Nitzsche and Welch | health_fitness | 41.28 | Ashley | Lopez    | F      | 9333 Valentine Point        | ... | 40.6729 | -73.5365  | 34496    | Librarian, public      | 1970-10-21 | c81755dbbbea |
| 3 | 3          | 2020-06-21 12:15:15   | 3591919803438423 | fraud_Haley Group                    | misc_pos       | 60.05 | Brian  | Williams | M      | 32941 Krystal Mill Apt. 552 | ... | 28.5697 | -80.8191  | 54767    | Set designer           | 1987-07-25 | 2159175b9e7  |
| 4 | 4          | 2020-06-21 12:15:17   | 3526826139003047 | fraud_Johnston-Casper                | travel         | 3.19  | Nathan | Massey   | M      | 5783 Evan Roads Apt. 465    | ... | 44.2529 | -85.0170  | 1126     | Furniture designer     | 1955-07-06 | 57ff021bd3f6 |

5 rows x 23 columns



2.) Select four columns to use as features (one just be trans\_date\_trans)

```
df_select = df[["trans_date_trans_time", "category", "amt", "city_pop", "is_fraud"]]

df_select.columns

Index(['trans_date_trans_time', 'category', 'amt', 'city_pop', 'is_fraud'], dtype='object')
```

3.) Create a your own variable out of trans\_date. Create dummies for factor vars

```
type(df_select["trans_date_trans_time"][0])

str

df_select["trans_date_trans_time"] = pd.to_datetime(df_select["trans_date_trans_time"])

<ipython-input-81-99f721e4ce0f>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
df_select["trans_date_trans_time"] = pd.to_datetime(df_select["trans_date_trans_time"])
```

```
dir(df_select["trans_date_trans_time"][0])

freq ,
'freqstr',
'fromisocalendar',
'fromisoformat',
'fromordinal',
'fromtimestamp',
'hour',
'is_leap_year',
'is_month_end',
'is_month_start',
'is_quarter_end',
'is_quarter_start',
'is_year_end',
'is_year_start',
'isocalendar',
'isoformat',
'isowekday',
'max',
'microsecond',
'min',
'minute',
'month',
'month_name',
'nanosecond',
'normalize',
'now',
'quarter',
'replace',
'resolution',
'round',
'second',
'strftime',
'strptime',
'time',
'timestamp'
```

```
timestamp',
'timetuple',
'timetz',
'to_datetime64',
'to_julian_date',
'to_numpy',
'to_period',
'to_pydatetime',
'today',
'toordinal',
'tz',
'tz_convert',
'tz_localize',
'tzinfo',
'tzname',
'utcfromtimestamp',
'utcnow',
'utcoffset',
'utctimetuple',
'value',
'week',
'weekday',
'weekofyear',
'year']
```

```
df_select["time_var"] = [i.second for i in df_select["trans_date_trans_time"]]
```

```
<ipython-input-83-fa4370ef92e9>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
df_select["time_var"] = [i.second for i in df_select["trans_date_trans_time"]]
```

```
X = pd.get_dummies(df_select, ["category"]).drop(["trans_date_trans_time", "is_fraud"], axis = 1)
y = df["is_fraud"]
```

X.head()

|   | amt   | city_pop | time_var | category_entertainment | category_food_dining | category_gas_transport | category_grocery_net | category_grocery_pos | category_ |
|---|-------|----------|----------|------------------------|----------------------|------------------------|----------------------|----------------------|-----------|
| 0 | 2.86  | 333497   | 25       | 0                      | 0                    | 0                      | 0                    | 0                    |           |
| 1 | 29.84 | 302      | 33       | 0                      | 0                    | 0                      | 0                    | 0                    |           |
| 2 | 41.28 | 34496    | 53       | 0                      | 0                    | 0                      | 0                    | 0                    |           |
| 3 | 60.05 | 54767    | 15       | 0                      | 0                    | 0                      | 0                    | 0                    |           |
| 4 | 3.19  | 1126     | 17       | 0                      | 0                    | 0                      | 0                    | 0                    |           |



➤ XXX SKIP THIS WE WILL TALK ABOUT NEXT CLASS

```
resample_X = X
resample_y = y
```

## ▼ 5.) Train a Logistic regression.

```
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
scaler = StandardScaler()
X_normalized = scaler.fit_transform(resample_X)
log_reg = LogisticRegression().fit(X_normalized, resample_y)
```

```
from sklearn.linear_model import LogisticRegression
```

```
log_reg = LogisticRegression().fit(X_normalized, resample_y)
```

## ▼ 6.) The company you are working for wants to target at a False Positive rate of 5% what threshold should you use? (Use oversampled data)

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
```

```
clf = LogisticRegression()
clf.fit(X_normalized, resample_y)
```

```
y_prob = clf.predict_proba(X_normalized)
```

```
target_fn_percentage = 5
threshold = np.percentile (y_prob[:,1], 100-target_fn_percentage)
```

```
y_pred = (y_prob[:,1] > threshold).astype(int)
```

```
confusion_matrix(resample_y, y_pred)
```

```
array([[527247,  26327],
       [   686,   1459]])
```

## ▼ 7.) If the company makes .02\*amt on True transactions and loses -amt on False (Use original data)


```
df_temp = df_select.copy()
```

```
df_temp["pred"] = log_reg.predict(resample_X)
```

```
/usr/local/lib/python3.8/dist-packages/sklearn/base.py:443: UserWarning: X has feature names, but LogisticRegression was fitted without feature names
warnings.warn(
```

```
df_temp = df_temp[["pred", "is_fraud", "amt"]]
```

```
df_temp.head()
```

|   | pred | is_fraud | amt   |  |
|---|------|----------|-------|-----------------------------------------------------------------------------------|
| 0 | 0    | 0        | 2.86  |                                                                                   |
| 1 | 0    | 0        | 29.84 |                                                                                   |
| 2 | 0    | 0        | 41.28 |                                                                                   |
| 3 | 0    | 0        | 60.05 |                                                                                   |
| 4 | 0    | 0        | 3.19  |                                                                                   |

```
df2 = df_temp.loc[df_temp['pred']==0,]
df2.dropna(axis=0, how='any', inplace=True)
df2 = df2.reset_index()
```

```
/usr/local/lib/python3.8/dist-packages/pandas/util/_decorators.py:311: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
return func(*args, **kwargs)
```

```
a= 0
for i in range(0,len(df2)):
    if df2.loc[i,'is_fraud']==0 :
        a = a + 0.02*df2.loc[i,'amt']
    if df2.loc[i,'is_fraud']==1 :
        a = a - df2.loc[i,'amt']
print("the profit is", a)
```

```
the profit is -35132.44080000058
```

▼ 8.) Using Logistic Regression Lasso to inform you. Would you use the selected features in a trusted prediction model?

```
LogisticRegression('l1')
```

```
LogisticRegression(penalty='l1')
```

```
# If most or all your variables go to 0 => Your data is garbage
# The regularization will tell us if our model has significance
```



```
# This of using coefficient strength similar to r^2

from sklearn.linear_model import LogisticRegressionCV
from sklearn.datasets import make_classification

clf = LogisticRegression(penalty='l1',solver="liblinear")
clf.fit(X_normalized, resample_y)
y_pred = clf.predict(X_normalized)
clf.coef_

array([[ 0.32966546, -0.12165055,  0.00675377, -0.06866738, -0.07648443,
         0.10641105,  0.0236613 ,  0.42938117, -0.06880143, -0.11334815,
        -0.09721065,  0.32442147,  0.          , -0.01000606,  0.40847022,
         0.1028411 , -0.91622471]])
```



We see that only 1 variable is 0. Therefore, I will use selected features in a trusted prediction model.

We see that only 1 variable is 0. Therefore, I will use selected features in a trusted prediction model.

Colab paid products - [Cancel contracts here](#)

✓ 40s    completed at 10:39 PM

